



*Go, change the world<sup>®</sup>*

# **ARTIFICIAL INTELLIGENCE IN AUTONOMOUS VEHICLES**

VII semester  
Department of AIML  
RV College of Engineering

Course Incharge: Dr. Viswavardhan Reddy Karna

**2024-2025**



# Course Contents – Unit - IV

*Go, change the world<sup>®</sup>*

## **Reinforcement Learning-Based Planning and Control:**

- Introduction to Reinforcement Learning.
- Learning Based Planning and Control in Autonomous Driving

## **Client Systems for Autonomous Driving:**

- Autonomous Driving: A Complex System
- Operating System for Autonomous Driving
- Computing Platform



# Reinforcement Learning-Based Planning and Control

*Go, change the world<sup>®</sup>*

## Key Characteristics of Reinforcement Learning

- Learning is interactive with the environment
- Closed-loop learning process
- The main entity: Agent
- Everything outside the agent: Environment

## Agent-Environment Interaction

- The agent makes decisions by taking actions
- The environment responds with new states and rewards
- Continuous iterative process



**Figure 7.1:** Reinforcement learning framework: agent interacting with environment by taking actions, sensing state, and getting rewards.



# Reinforcement Learning-Based Planning and Control

*Go, change the world<sup>®</sup>*

## The Learning Process

- Learning occurs in rounds indexed by time ( $t = 0, 1, 2, 3 \dots$ )
- At each time step  $t$ :
  - The agent perceives the environment (State  $S_t$ )
  - Takes an action ( $A_t$ )
  - Receives a new state ( $S_{t+1}$ ) and a reward ( $R_t$ )

## Reward Function

- Defines how rewards are generated by the environment
- Maps state or state-action pairs to a scalar reward
- The agent aims to maximize cumulative rewards



# Reinforcement Learning-Based Planning and Control

*Go, change the world<sup>®</sup>*

## Episodic vs. Continuing Tasks

### • Episodic Learning:

- Fixed sequences of states
- Ends in terminal states

### • Continuing Tasks:

- Infinite sequence of states
- Requires discounted reward computation

## Return Computation

- Episodic tasks: Sum of rewards
- Continuing tasks: Discounted rewards
- Return function:

- $G_t = R_{t+1} + \gamma R_{t+2} + \gamma^2 R_{t+3} \dots$
- $\gamma$  (gamma) is the discount factor ( $0 \leq \gamma \leq 1$ )

$$G_t = R_{t+1} + \gamma R_{t+2} + \gamma^2 R_{t+3} \dots = \sum_{k=0}^{\infty} \gamma^k R_{t+k+1},$$



# Reinforcement Learning-Based Planning and Control

*Go, change the world<sup>®</sup>*

## Policy and Decision Making

- Policy ( $\pi$ ): Strategy for decision making
- Maps states to actions
- Can be:
  - Simple (lookup table)
  - Complex (deep neural network)

## Value Functions

- Measures the expected return for being in a state
- State-Value Function:  $V_{\pi}(s) = E_{\pi} (G_t | S_t = s)$
- State-Action Value Function (Q-value):  $Q_{\pi}(s,a)$



# Reinforcement Learning-Based Planning and Control

*Go, change the world<sup>®</sup>*

in Chapter 6) where the state transitions are Markov, the value function could be written as:  $V_{\pi}(s) = E_{\pi} (G_t \mid S_t = s) = E_{\pi} (\sum_{k=0}^{\infty} \gamma^k R_{t+k+1} \mid S_t = s)$ , where  $E_{\pi}$  is the expectation operator given the policy  $\pi$ . This value function is also called the *state-value function for policy  $\pi$* . Similarly, the Q-value function is the mapping of state-action pairs to a scalar value representing the expected return at state  $S_t$  after taking action  $A_t$  and following the policy  $\pi$  afterward. It is denoted as:  $Q_{\pi}(s, a) = E_{\pi} (G_t \mid S_t = s, A_t = a) = E_{\pi} (\sum_{k=0}^{\infty} \gamma^k R_{t+k+1} \mid S_t = s, A_t = a)$ .



# Reinforcement Learning-Based Planning and Control

*Go, change the world<sup>®</sup>*

- **Q-Learning** is one of the most popular algorithms used for reinforcement learning.
- The idea is to learn to approximate the strategy that maximizes the expected return at any given state  $s_t$  by taking action  $a_t$  and then following the optimal strategy:  $Q(s_t, a_t) = \max_{\pi} R_{t+1}$ .
- How we then select the policy is based on:  $\pi(s_t) = \operatorname{argmax}_{a_t} Q(s_t, a_t)$ .
- The key problem in Q-learning is how to accurately estimate the Q-function that maps state-action pairs to an optimal expected return





# Reinforcement Learning-Based Planning and Control

*Go, change the world<sup>®</sup>*

- **Q-Learning** is one of the most popular algorithms used for reinforcement learning.
- The idea is to learn to approximate the strategy that maximizes the expected return at any given state  $s_t$  by taking action  $a_t$  and then following the optimal strategy:  $Q(s_t, a_t) = \max_{\pi} R_{t+1}$ .
- How we then select the policy is based on:  $\pi(s_t) = \operatorname{argmax}_{a_t} Q(s_t, a_t)$ .
- The key problem in Q-learning is how to accurately estimate the Q-function that maps state-action pairs to an optimal expected return



# Reinforcement Learning-Based Planning and Control

## Q-Learning algorithm

*Go, change the world<sup>®</sup>*

```
1 function Q_Learning(Episodes)
2   Initialize the  $Q(s, a)$ , and  $Q(\text{terminate-state}, :) = 0$ 
3   for each episode in Episodes:
4     Initialize to start state  $S$ 
5     repeat (for each step of episode):
6       choose an action  $a \in A$  using the policy from the  $Q$  table with  $\epsilon$ -greedy
7       Take the action  $a$  and observe the reward  $R$  and the next state  $s'$ 
8        $Q(s, a) \leftarrow Q(s, a) + \alpha[R + \gamma \max_{a'} Q(s', a') - Q(s, a)]$ 
9        $S \leftarrow s'$ 
10    until  $S$  is a terminal state
```



# Reinforcement Learning-Based Planning and Control

## Deep-Q-Learning algorithm.

*Go, change the world<sup>®</sup>*

```
1 function DQN_Learning(Episodes)
2   Initialize the  $Q(s,a)$  with random weights
3   Initialize the replay memory  $D$ 
4   for each episode in Episodes:
5     repeat :
6       choose an action  $a \in A$  using the policy from the  $Q$  table with  $\epsilon$ -greedy
7       take the action  $a$  and observe the reward  $r$  and the next state  $s'$ 
8       store experience  $(s, a, r, s')$  in replay memory  $D$ 
9       sample random transition  $(ss, aa, rr, ss')$  from replay memory  $D$ 
10      calculate target for each minibatch transition
11        if is terminal state:
12           $tt = rr$ 
13        else:
14           $tt = rr + \gamma \max_{a'} Q(s', a') - Q(s, a)$ 
15        train deep neural network  $Q$  with  $Loss = \frac{1}{2} (tt - Q(ss, aa))^2$ 
16       $s \leftarrow s'$ 
17    until  $s$  is a terminal state
```

The neural network  $Q(s, a)$  is trained using a squared loss function:

$$Loss = \frac{1}{2} (tt - Q(ss, aa))^2$$

This updates the network weights to make  $Q(s, a)$  closer to  $tt$ .



# Reinforcement Learning-Based Planning and Control

## Deep-Q-Learning algorithm.

*Go, change the world<sup>®</sup>*

```
1 function DQN_Learning(Episodes)
2   Initialize the  $Q(s,a)$  with random weights
3   Initialize the replay memory  $D$ 
4   for each episode in Episodes:
5     repeat :
6       choose an action  $a \in A$  using the policy from the  $Q$  table with  $\epsilon$ -greedy
7       take the action  $a$  and observe the reward  $r$  and the next state  $s'$ 
8       store experience  $(s, a, r, s')$  in replay memory  $D$ 
9       sample random transition  $(ss, aa, rr, ss')$  from replay memory  $D$ 
10      calculate target for each minibatch transition
11        if is terminal state:
12           $tt = rr$ 
13        else:
14           $tt = rr + \gamma \max_{a'} Q(s', a') - Q(s, a)$ 
15        train deep neural network  $Q$  with  $Loss = 1/2 (tt - Q(ss, aa))^2$ 
16       $s \leftarrow s'$ 
17    until  $s$  is a terminal state
```

- If  $s'$  is a terminal state, the target Q-value  $tt$  is simply the observed reward ( $rr$ ).

- Otherwise, the Bellman equation is used:

$$tt = r + \gamma \max_{a'} Q(s', a')$$

- $\gamma$  (gamma) is the discount factor ( $0 \leq \gamma \leq 1$ ).
- $\max Q(s', a')$  is the highest predicted Q-value for the next state.
- The difference  $tt - Q(s, a)$  is the TD (Temporal Difference) error.

# Reinforcement Learning-Based Planning and Control

## Deep-Q-Learning algorithm.

*Go, change the world<sup>®</sup>*

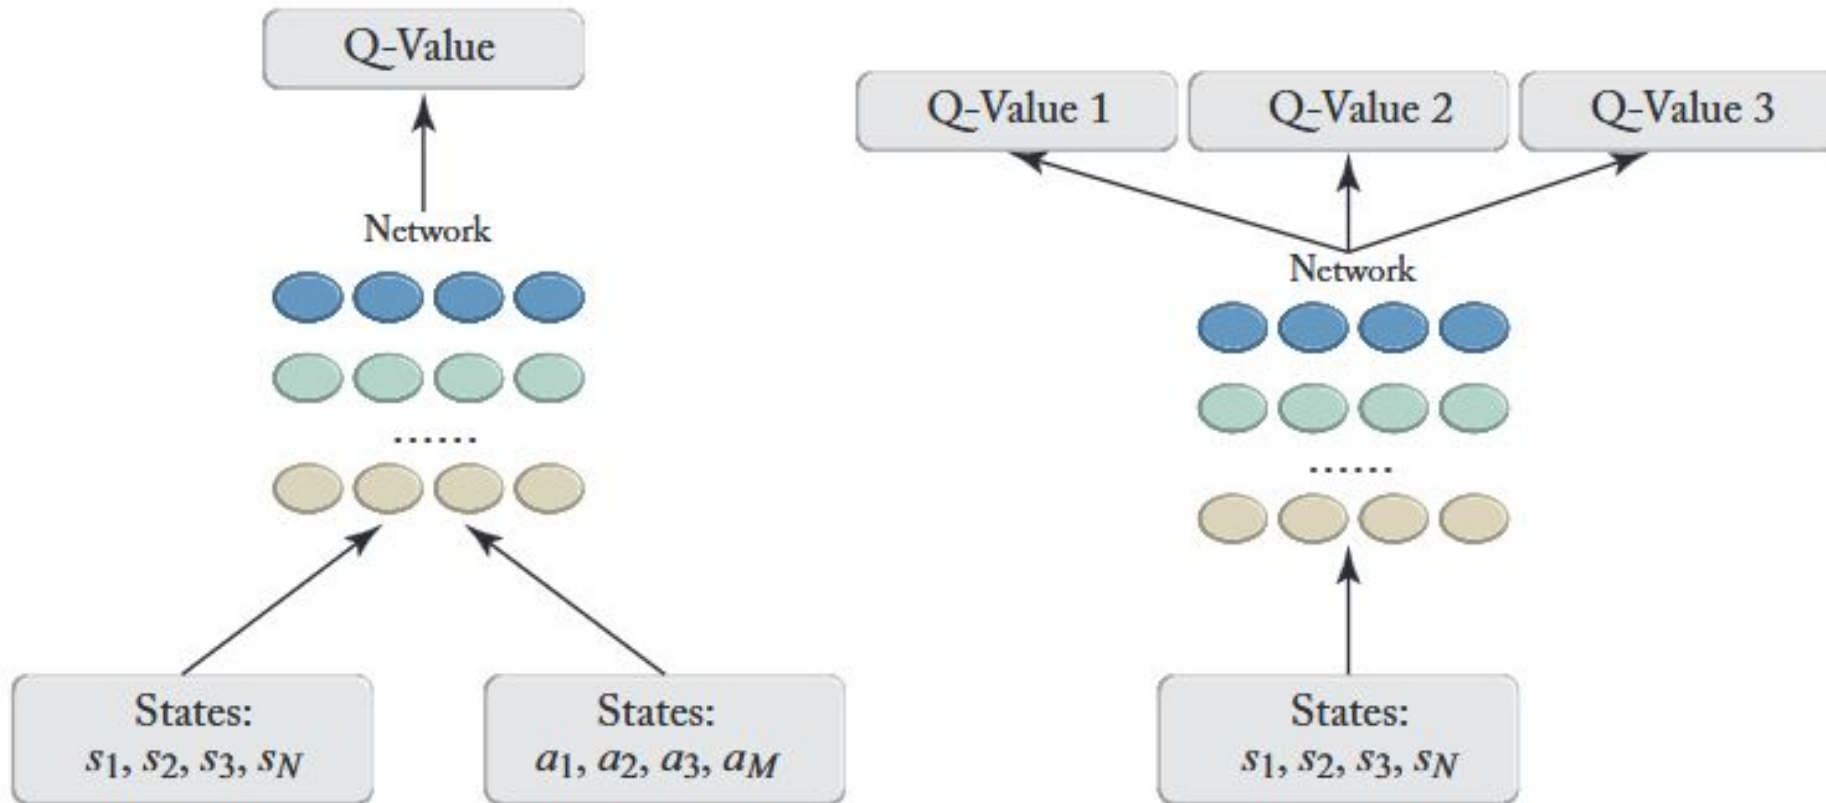


Figure 7.3: Deep-Q-Learning (DQN): Neural network structures.



# Reinforcement Learning-Based Planning and Control

*Go, change the world<sup>®</sup>*

## ACTOR-CRITIC METHODS

- Asynchronous Advantage Actor-Critic (A3C) algorithm presented by Google's DeepMind.
- It has been shown to be faster, simpler, and more robust than the traditional Deep-Q-Learning algorithms on standard reinforcement learning tasks
- The structure of an A3C network differs from DQN by using multiple learning agents instead of one.
- Each agent has its own neural network and interacts with its own environment.
- A global network aggregates updates from worker agents.
- Asynchronous learning: Each agent updates its policies independently before synchronizing with the global network.





# Reinforcement Learning-Based Planning and Control

## A3C Training Workflow

*Go, change the world<sup>®</sup>*

- .Initialize a global network and reset each individual worker network to the global network.
- .Each individual worker agent interacts and learns within its own environment.
- .Each individual worker agent computes the loss function for its neural network.
- .Each individual worker agent updates its gradient from the computed loss.
- .Worker agents together update the global network with the appropriate gradients.

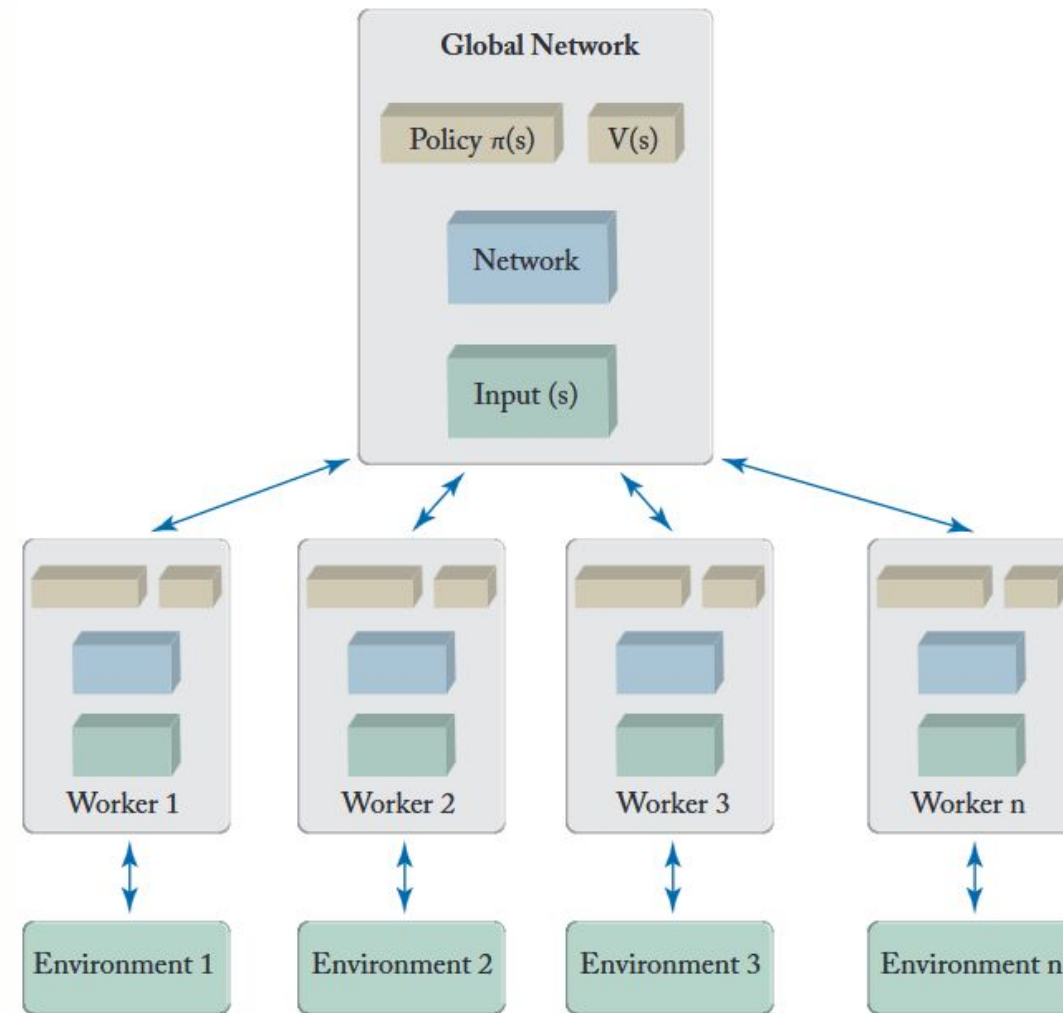


Figure 7.6: Asynchronous Advantage Actor-Critic Framework (based on [11]).



# Reinforcement Learning-Based Planning and Control

*Go, change the world<sup>®</sup>*

## Shared Network Architecture

- Both value and policy networks share:
  - **Convolutional Layers** for spatial feature extraction.
  - **LSTM Layer** for handling temporal dependencies.
- Helps maintain location invariance and sequence learning.

## Actor-Critic Mechanism

- **Critic (Value Function  $V\pi(s)$ ):**
  - Evaluates the quality of a state.
  - Guides policy updates.
- **Actor (Policy  $\pi(s)$ ):**
  - Determines the optimal action.
  - Adjusts based on critic feedback.

## Advantage Function

- Measures improvement over expected return:
  - $A(s, a) = Q(s, a) - V(s)$
  - Helps refine policy updates.
- Advantage function ensures more effective training.





# Reinforcement Learning-Based Planning and Control

*Go, change the world<sup>®</sup>*

## Loss Functions

- **Value Loss:**

- $L_{\text{value}} = \sum (R - V(s))^2$

- **Policy Loss:**

- $L_{\text{policy}} = A(s) * \log(\pi(s)) + H\pi * \beta$

- Includes entropy term  $H\pi$  for balancing exploration and exploitation.

## Exploration vs. Exploitation

- **Entropy ( $H\pi$ ) measures policy uncertainty.**

- High entropy  $\rightarrow$  More exploration.
  - Low entropy  $\rightarrow$  More exploitation.

- Helps stabilize learning and avoid local optima.



# LEARNING-BASED PLANNING AND CONTROL IN AUTONOMOUS DRIVING

*Go, change the world<sup>®</sup>*

## REINFORCEMENT LEARNING ON BEHAVIORAL DECISION

- The main goal of applying reinforcement learning (RL) in behavioral decisions is to **manage diverse traffic scenarios**.
- **Strict adherence to traffic rules** is not always effective in real-world driving.
- **"Long-tail"** cases in behavioral decisions require adaptive decision-making.
- **Human driving experiences** can serve as valuable **examples** for **RL-based systems**.
- RL can complement rule-based behavioral decision approaches, which remain the industry standard.



# LEARNING-BASED PLANNING AND CONTROL IN AUTONOMOUS DRIVING

*Go, change the world<sup>®</sup>*

## REINFORCEMENT LEARNING ON BEHAVIORAL DECISION

### RL-Based Behavioral Decision Model

- The action space (Desires) is defined as:

$$\mathbf{D} = [0, v_{\max}] \times \mathbf{L} \times \{\mathbf{g}, \mathbf{t}, \mathbf{o}\}^n$$

where:

- $v_{\max}$ : Desired target speed of the vehicle
- $\mathbf{L}$ : Set of discrete lateral lane positions
- $\mathbf{g}$ : Give way (yield),      $\mathbf{t}$ : Take (overtake)
- $\mathbf{o}$ : Keep an offset distance (nudge/attention)



# LEARNING-BASED PLANNING AND CONTROL IN AUTONOMOUS DRIVING

*Go, change the world<sup>®</sup>*

## REINFORCEMENT LEARNING ON BEHAVIORAL DECISION

### State Space

- The environment model is generated by interpreting sensory information.
- Includes additional information such as kinematics of moving objects.

### Implementation & Limitations

- RL agent initialized with imitation learning.
- Updated using iterative Policy Gradient approach.



# LEARNING-BASED PLANNING AND CONTROL IN AUTONOMOUS DRIVING

*Go, change the world<sup>®</sup>*

## REINFORCEMENT LEARNING ON PLANNING AND CONTROL

- The key challenge in RL-based planning and control is **how to design the state spaces**.
- To compute the **motion planning or feedback control level actions**, it is necessary to include both **autonomous vehicle information** and the **surrounding environment**.
- If we won't take sensory data – which is raw, the state spaces will somehow incorporate structured information about AV and its surroundings - **large multi-dimensional continuous space**.
- To tackle the challenge of a continuous state space, **cell-mapping techniques** can be combined with **reinforcement learning**.



# LEARNING-BASED PLANNING AND CONTROL IN AUTONOMOUS DRIVING

*Go, change the world<sup>®</sup>*

## REINFORCEMENT LEARNING ON PLANNING AND CONTROL

### Continuous State Space

- In reinforcement learning (RL), the **state space** represents all possible situations an agent can encounter.
- A **continuous state space** means the states are not discrete but can take on any value within a given range.
- Example: In autonomous driving, the **position and speed of a car** can take infinitely many values rather than being confined to a limited set of predefined states.

### Cell-Mapping Techniques

- **Cell-mapping** is a method that divides a **continuous state space into discrete cells**. Each cell represents a group of **similar states**



# LEARNING-BASED PLANNING AND CONTROL IN AUTONOMOUS DRIVING

*Go, change the world<sup>®</sup>*

## REINFORCEMENT LEARNING ON PLANNING AND CONTROL

### Combining Cell-Mapping with RL

- By **discretizing the state space** into manageable cells, RL algorithms can operate more effectively without requiring **infinite memory or computational power**.
- This hybrid approach is useful for problems where a fully continuous representation is too complex but full discretization is too simplistic.
- It can improve convergence speed and stability in RL applications like robotic control, traffic management, and autonomous systems.



# LEARNING-BASED PLANNING AND CONTROL IN AUTONOMOUS DRIVING

*Go, change the world<sup>®</sup>*

## REINFORCEMENT LEARNING ON PLANNING AND CONTROL

State Symbol	State Variable	Range
$X_1 = v$	Velocity	$-1.5 \leq X_1 \leq 1.5$ "m/s"
$X_2 = x$	X Cartesian Coordinate	$-0.9 \leq X_2 \leq 0.9$ "m"
$X_2 = y$	Y Cartesian Coordinate	$-1.3 \leq X_3 \leq 1.3$ "m"
$X_2 = \theta$	Orientation	$-\pi \leq X_4 \leq \pi$ "rad"

Action Symbol	Action Values
Voltage in traction motor	-18 V   0 V   18 V
Steering angle	-23°   0°   23°

### Algorithm 1 CACM-RL

```
1:  Initialise  $Q\text{-Table}(s,a)$  y  $Model\_Table$ 
2:   $x \leftarrow$  current state
3:   $s \leftarrow cell(x)$ 
4:  IF       $s \in \text{drain}$  or  $s \in \text{goal}$  or  $s \in \text{safety\_area}$ 
5:  THEN  $F_{\text{reactive}}(x)$ 
6:  ELSE IF       $D\text{-}k\text{-adjoining}(x,x')$ 
7:  THEN  $Q\text{-Table}(s,a) \leftarrow s',r$ 
8:       $Model\_Table \leftarrow IT(x,x')$ 
9:       $a \leftarrow policy(s)$ 
10:     Execute action  $a$  on the vehicle
11:     Observe the new state  $x'$  and  $r$ 
12: UNTIL the end of the learning stage
13: FOR all  $(s,a)$ , repeat  $N$  times
14:      $\bar{x}' \leftarrow Model\_Table, DT(x,x')$ 
15:      $\bar{s}' \leftarrow cell(\bar{x}')$ 
16:      $Q\text{-Table}(s,a) \leftarrow \bar{s}',r$ 
```





# LEARNING-BASED PLANNING AND CONTROL IN AUTONOMOUS DRIVING

*Go, change the world<sup>®</sup>*

## REINFORCEMENT LEARNING ON PLANNING AND CONTROL

### Initialization Phase

#### 1. Initialize Q-Table(s, a) and Model\_Table

- The **Q-Table** stores the expected rewards for taking actions in different states.
- The **Model\_Table** keeps track of the transitions and rewards for each state-action pair.

#### 2. $x \leftarrow$ current state

- The current state of the system (e.g., velocity, position, orientation) is assigned to  $x$ .

#### 3. $s \leftarrow$ cell( $x$ )

- The continuous state  $x$  is mapped to a **cell** in the discretized state space. This step is crucial for handling continuous environments in reinforcement learning.



# LEARNING-BASED PLANNING AND CONTROL IN AUTONOMOUS DRIVING

*Go, change the world<sup>®</sup>*

## REINFORCEMENT LEARNING ON PLANNING AND CONTROL

### State Classification & Action Selection Phase

4. IF  $s$  is drain OR  $s$  is goal OR  $s$  is safety\_area

- If the agent reaches a **drain state** (undesirable), a **goal state** (successful outcome), or a **safety area**, then:

5. THEN  $F_{\text{reactive}}(x)$

- Apply a reactive function to handle these special cases. This function could reset the agent or enforce constraints.

6. ELSE IF  $D\text{-}k\text{-}adjoining(x, x')$

- Check if the new state  $x'$  is **D-k adjoining**, meaning it is within a defined neighborhood of the current state  $x$ .

7. THEN  $Q\text{-}Table(s, a) \leftarrow s', r$

- Update the Q-Table by storing the new state  $s'$  and the reward  $r$  for taking action  $a$  in state  $s$ .

8.  $Model\_Table \leftarrow (x, \pi(x'))$

- Update the **Model\_Table** with the new state transition  $(x, \pi(x'))$ , where  $\pi(x')$  represents a transition function that determines the next state.

9.  $a \leftarrow policy(s)$

- Select an action  $a$  based on the policy, which could be an  **$\epsilon$ -greedy policy**, **softmax**, or another decision-making strategy.



# LEARNING-BASED PLANNING AND CONTROL IN AUTONOMOUS DRIVING

*Go, change the world<sup>®</sup>*

## REINFORCEMENT LEARNING ON PLANNING AND CONTROL

### Action Execution & Learning Phase

#### 10. Execute action $a$ on the vehicle

- Apply the chosen action  $a$  to the system (e.g., adjust voltage or steering angle).

#### 11. Observe the new state $x'$ and $r$

- The agent receives feedback by observing the new state  $x'$  and the reward  $r$  from the environment.

#### 12. UNTIL the end of the learning stage

- Repeat the process until the training phase is complete.

#### 13. FOR ALL $(s, a)$ , repeat $N$ times

- Perform updates for all state-action pairs multiple times to refine learning.



# LEARNING-BASED PLANNING AND CONTROL IN AUTONOMOUS DRIVING

*Go, change the world<sup>®</sup>*

## REINFORCEMENT LEARNING ON PLANNING AND CONTROL

- $\mathbf{x}' \leftarrow \text{Model\_Table}, \text{DT}(\mathbf{x}', \mathbf{x}')$
- Retrieve the next predicted state from the **Model\_Table** using a **Dynamic Transition (DT)** function.
- $\mathbf{s}' \leftarrow \text{cell}(\mathbf{x}')$
- Convert the continuous state  $\mathbf{x}'$  into its corresponding cell representation.
- $\mathbf{Q\text{-}Table}(\mathbf{s}, \mathbf{a}) \leftarrow \mathbf{g}', \mathbf{r}'$
- Final Q-Table update with the reward  $\mathbf{r}'$  and new state  $\mathbf{g}'$ .



# Client Systems for Autonomous Driving

*Go, change the world<sup>®</sup>*

- Autonomous driving is an integration of multiple technologies.
- The client system consists of:
  - **Operating system and Hardware platform**
- Coordination between these components is essential for performance.

## Hardware Platform Overview

- Follows the **sensing-perception-action** paradigm.
- **Sensors** collect environmental data.
- **Computing platform** processes data for perception and action computation.
- **Control platform** executes the action plans.

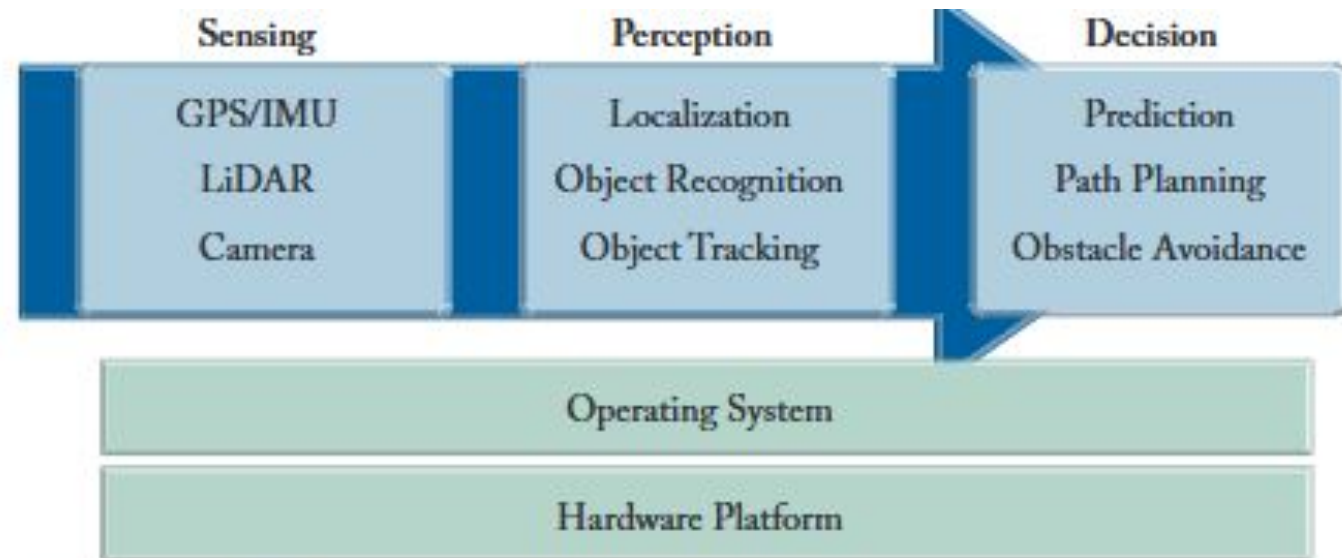


Figure 8.1: Hardware platform for autonomous driving.



# Client Systems for Autonomous Driving

*Go, change the world<sup>®</sup>*

## Role of the Operating System

- Manages communication between hardware components.
- Coordinates **resource allocation** for real-time tasks.

## Real-Time Processing Requirements

- Example: **Camera requires 60 frames per second.**
- Each frame must be processed in **less than 16 ms.**
- Increased data load affects processing efficiency.

## Challenges in Resource Allocation

- Data bursts** (e.g., LiDAR point clouds) can overload CPU.
- Can lead to **dropped frames** from the camera.
- Need a mechanism to restrict resource usage per component.

# Client Systems for Autonomous Driving

*Go, change the world<sup>®</sup>*

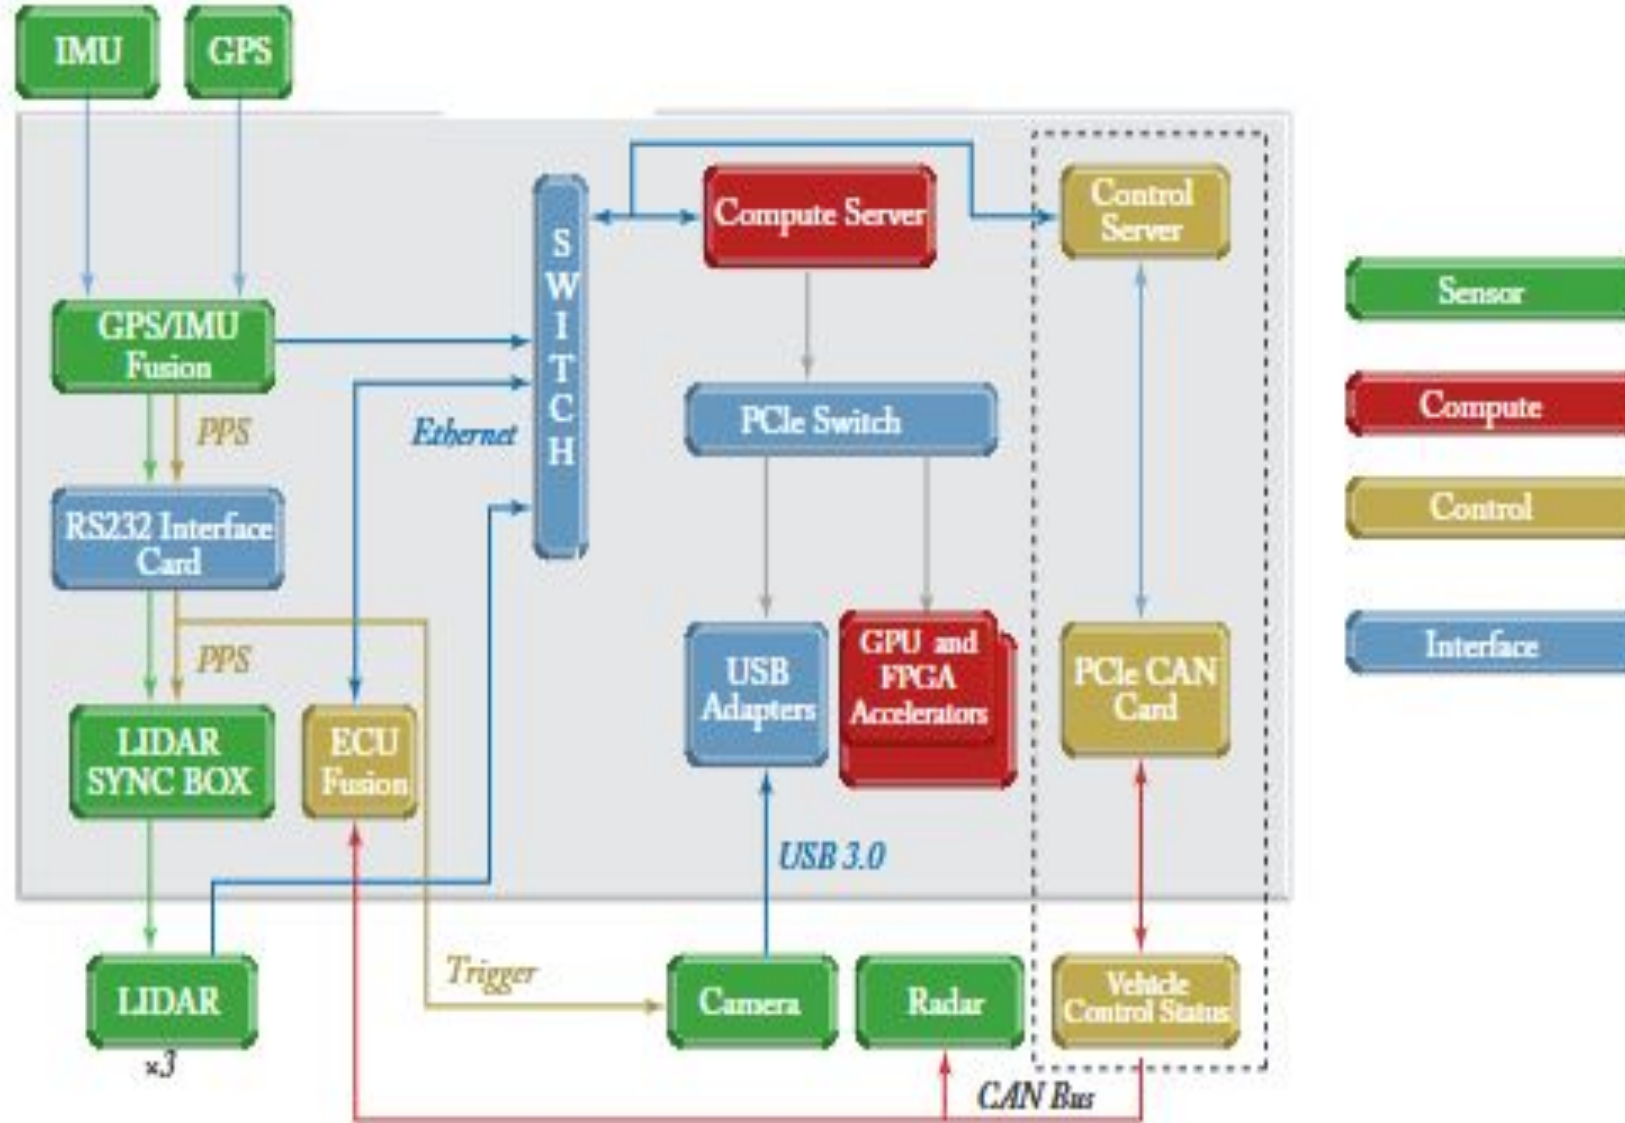


Figure 8.2: Hardware platform for autonomous driving.





# Client Systems for Autonomous Driving

*Go, change the world<sup>®</sup>*

- The two main functions provided by the operating system include communication, and resource allocation.
- The main components are divided into 3 types: **ROS Master**, **ROS Node**, and **ROS Service**.
- Function of ROS master is to provide name service.
- It stores the operating parameters that are required at **startup**, the **name of the connection between the upstream node and the downstream node**, and the **name of the existing ROS services**.

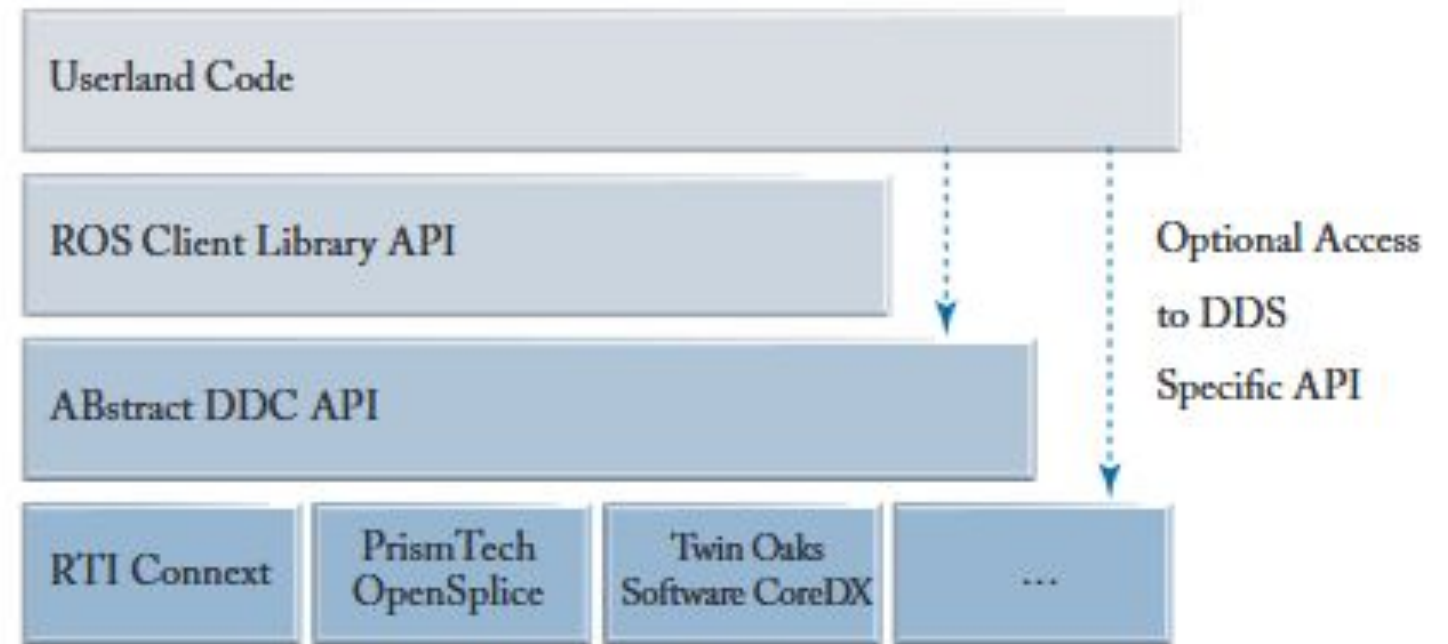


Figure 8.3: ROS 2.0 communication with DDS.



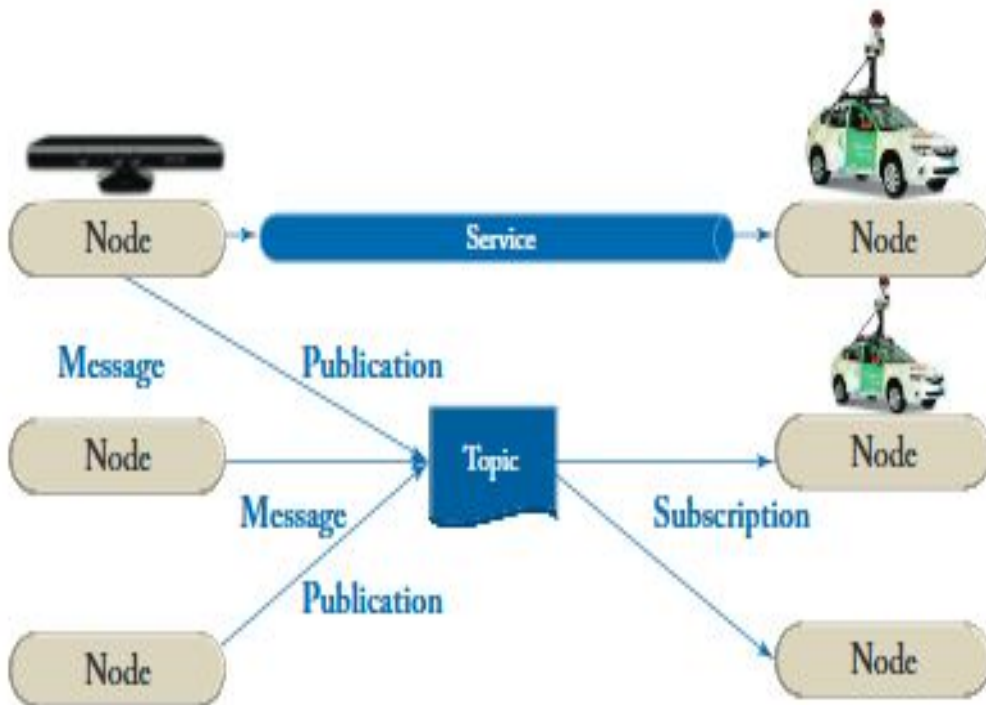
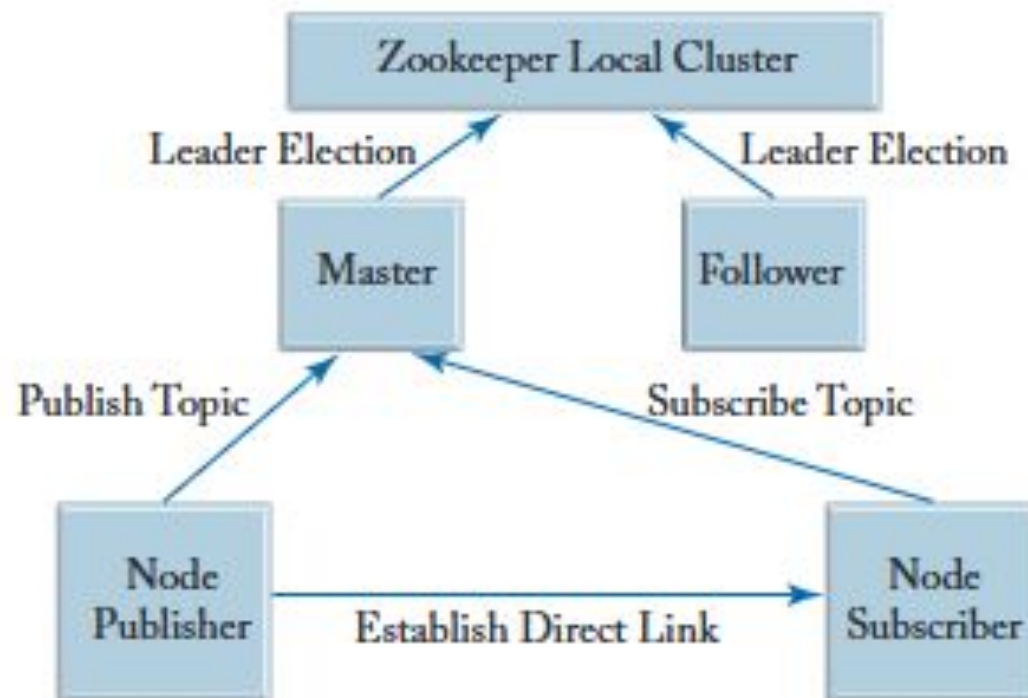


Figure 8.4: ROS communication mechanisms.



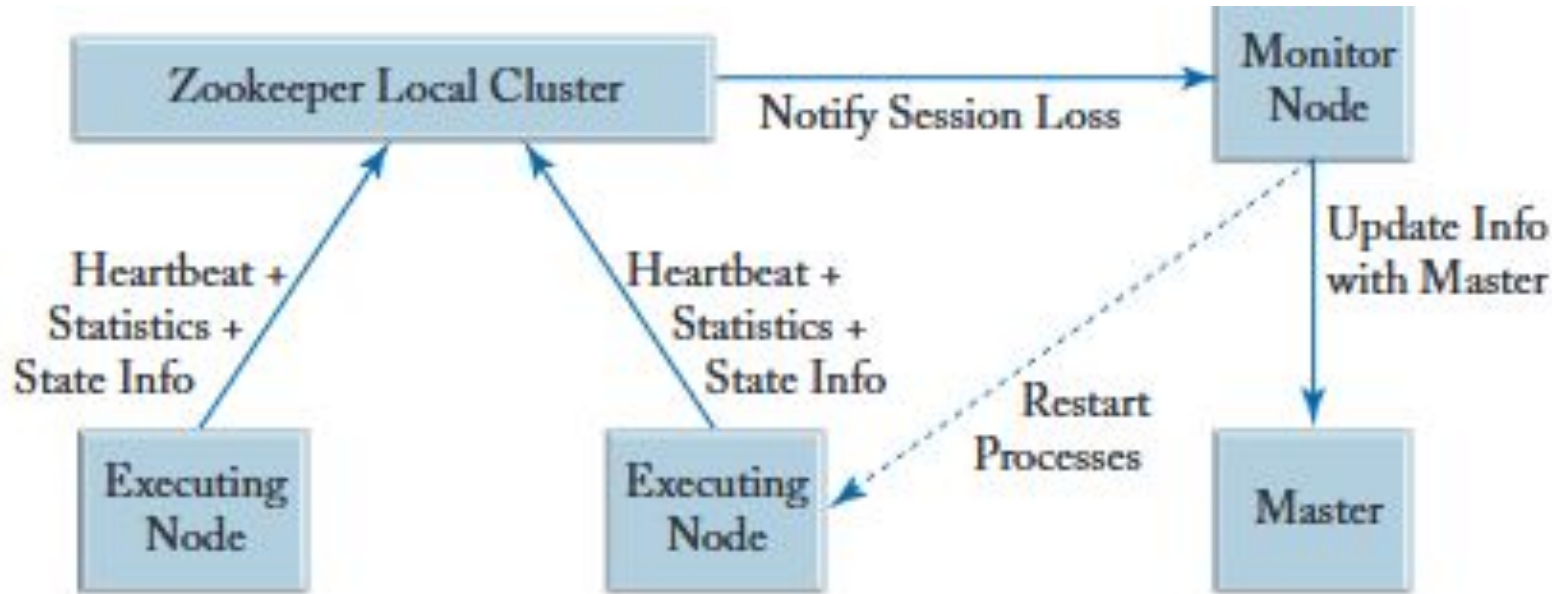


Figure 8.6: Monitor node.

- **system reliability:** If an autonomous vehicle is moving on the road, suddenly the **ROS master node crashes**, leading to a **system shutdown**.
- one master in the whole system, but it is not acceptable in autonomous vehicle applications.
- Thus we need to decentralize the master and achieve robustness and reliability - Zookeeper



# Client Systems for Autonomous Driving - Performance

*Go, change the world<sup>®</sup>*

- Un-necessary delay
- Memory consumption
- Waste of computing and memory resources – due to broadcast



# Client Systems for Autonomous Driving - COMPUTING PLATFORM

*Go, change the world<sup>®</sup>*

## Existing Autonomous Driving Computing Platform

### •Hardware Configuration:

- Two compute boxes with Intel Xeon E5 processors
- 4-8 Nvidia K80 GPUs per box
- Connected via PCI-E bus

### •Performance Metrics:

- 12-core CPU delivers 400 GOPS/s
- Each GPU: 8 TOPS/s, 300 W power consumption
- Overall: 64.5 TOPS/s, ~3,000 W total power

### •Challenges:

- High power consumption (~5,000 W max)
- High cost (\$20,000–\$30,000 per box)

### •Existing Computing Solutions

#### •Overview of major computing architectures:

- GPU-based solutions
- DSP-based solutions
- FPGA-based solutions
- ASIC-based solutions



# Client Systems for Autonomous Driving - COMPUTING PLATFORM

*Go, change the world<sup>®</sup>*

## GPU-Based Computing Solutions

### •Nvidia PX Platform

- 2 Tegra SoCs + 2 Pascal GPUs
- Dedicated memory & specialized DNN instructions
- PCI-E Gen 2  $\times$  4 bus (4.0 GB/s)
- Gigabit Ethernet interconnection (70 GB/s)
- 24 TOPS/s performance
- 2,800 images/s on AlexNet

## DSP-Based Solutions

### •Texas Instruments TDA2x SoC

- 2 floating-point C66x DSP cores
- 4 Vision Accelerators (8 $\times$  acceleration vs. Cortex-A15 CPU)

### •CEVA XM4 DSP

- Optimized for computer vision tasks
- Energy-efficient: <30 mW for 1080p at 30 FPS



# Client Systems for Autonomous Driving - COMPUTING PLATFORM

*Go, change the world<sup>®</sup>*

## FPGA-Based Solutions

- **Altera Cyclone V SoC** (Used in Audi vehicles)
  - Optimized for sensor fusion & object detection
- **Xilinx Zynq UltraScale MPSoC**
  - CNN processing: 14 images/sec/Watt
  - Object tracking: 60 FPS for 1080p live stream

## ASIC-Based Solutions

- **MobilEye EyeQ5**
  - Heterogeneous, fully programmable accelerators
  - Optimized for computer vision, signal processing, ML tasks
  - Multiple EyeQ5 devices interconnected via PCI-E



# Client Systems for Autonomous Driving - COMPUTING PLATFORM

*Go, change the world<sup>®</sup>*

## Computer Architecture Design Exploration

### •Key research questions:

- Best computing units for specific workloads
- Feasibility of mobile processors for autonomous driving
- Efficient platform design strategies

## Matching Workloads to Computing Units

### •Convolution & Feature Extraction Workloads:

- CPU: 8 ms/convolution, 20 mJ
- GPU: 2 ms/convolution, 4.5 mJ (most efficient for convolution)
- DSP: 5 ms/convolution, 7.5 mJ

### •Feature Extraction Task:

- CPU: 20 ms, 50 mJ
- GPU: 10 ms, 22.5 mJ
- DSP: 4 ms, 6 mJ (most efficient for feature extraction)

### •Conclusion:

- GPUs excel at convolution tasks
- DSPs excel at feature processing
- CPUs handle control-heavy tasks better

# Client Systems for Autonomous Driving

*Go, change the world<sup>®</sup>*

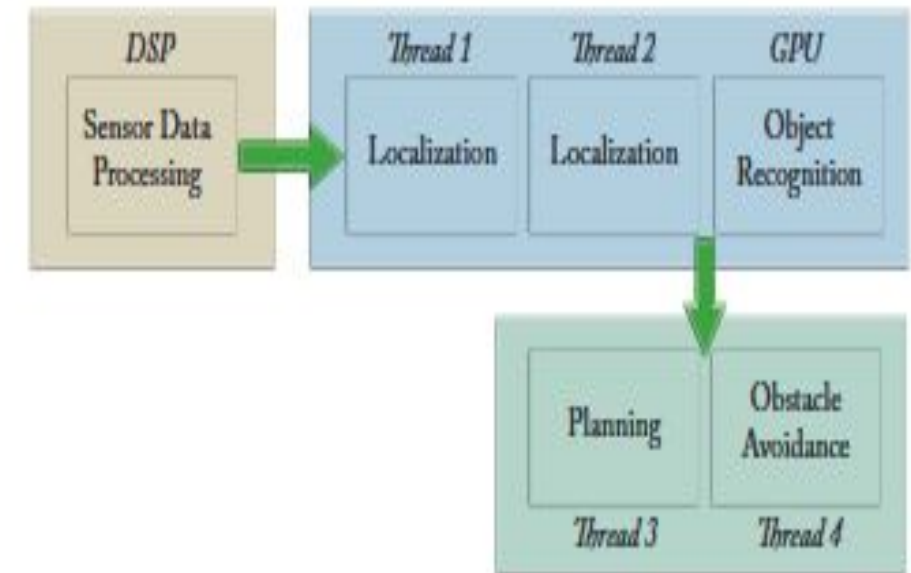



Figure 8.8: Autonomous navigation system on mobile SoC.





# Client Systems for Autonomous Driving

*Go, change the world<sup>®</sup>*

## ARM Mobile SoC Performance

### •Localization Pipeline:

- Processes 25 images/sec (almost keeping up with 30 images/sec generation)

### •Deep Learning Pipeline:

- 2-3 object recognition tasks per second

### •Planning & Control Pipeline:

- Plans a path within 6 ms

### •Power Consumption:

- Average consumption: 11 W

### •Performance Observations:

- Vehicle maintained localization at 5 mph
- With more computing resources, higher speeds could be achieved
- Potential for a production-level autonomous driving system



# Client Systems for Autonomous Driving

*Go, change the world<sup>®</sup>*

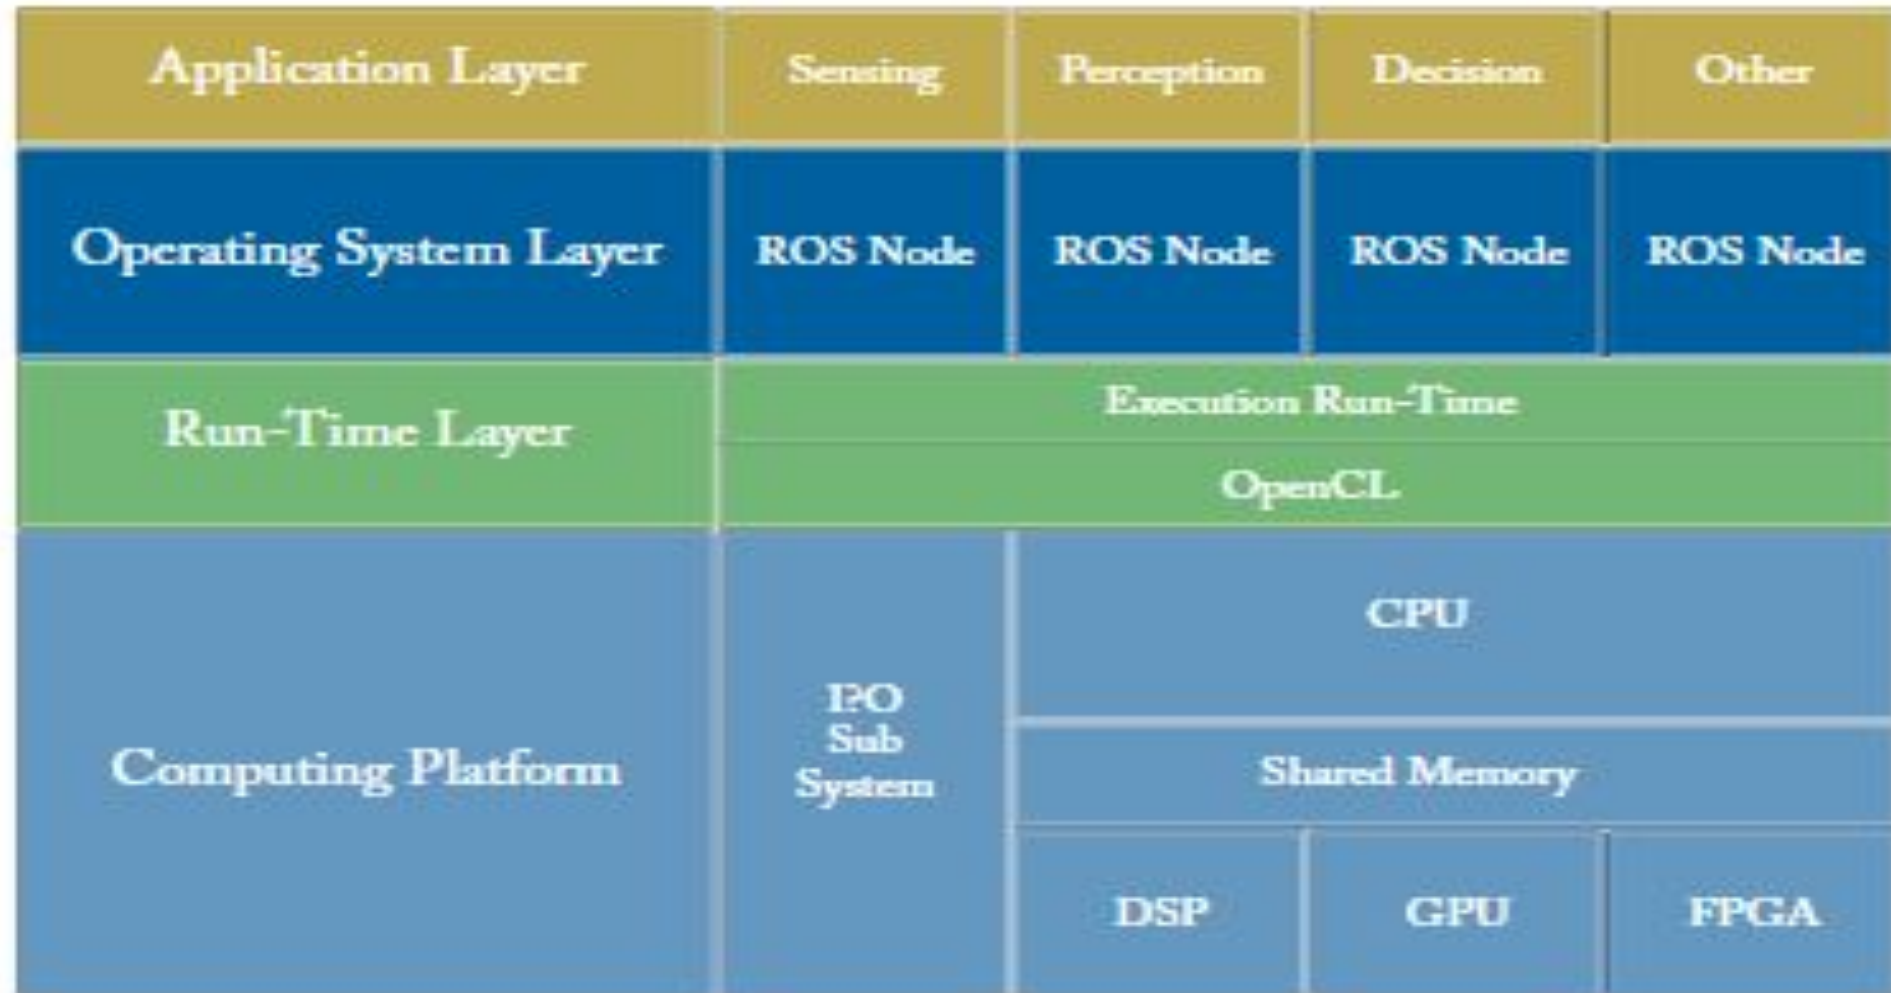


Figure 8.9: Computing stack for autonomous driving.



# Client Systems for Autonomous Driving

*Go, change the world<sup>®</sup>*

- **Application Layer (Top - Yellow)**

- Handles high-level tasks such as:

- **Sensing:** Collecting data from cameras, LiDAR, radar, and other sensors.
- **Perception:** Processing sensory data for object detection, localization, and mapping.
- **Decision:** Making driving decisions (e.g., lane changes, braking).
- **Other:** Additional functionalities like communication with external systems.

- **Operating System Layer (Dark Blue)**

- Composed of **ROS (Robot Operating System) Nodes**, which manage different software modules required for autonomous driving.

- Each **ROS Node** handles specific tasks like perception, decision-making, or sensor data processing.



# Client Systems for Autonomous Driving

*Go, change the world<sup>®</sup>*

- **Run-Time Layer (Green)**

- Includes **Execution Run-Time** and **OpenCL**.
- **Execution Run-Time** manages the execution of tasks in real-time.
- **OpenCL** provides a framework for running applications on heterogeneous hardware (CPU, GPU, FPGA).

- **Computing Platform (Bottom - Blue)**

- The hardware foundation for autonomous driving computations.
  - **CPU (Central Processing Unit)**: Handles general computing tasks.
  - **GPU (Graphics Processing Unit)**: Optimized for deep learning and parallel processing tasks.
  - **DSP (Digital Signal Processor)**: Specialized for low-power, high-speed signal processing.
  - **FPGA (Field-Programmable Gate Array)**: Provides flexible, energy-efficient hardware acceleration.
  - **Shared Memory**: Enables data sharing between different computing units.
  - **I/O Subsystem**: Manages data transfer between sensors and the computing platform.